

HiLabs Workshop: Reliability & Evaluation Report

1. Quantitative Evaluation Summary

The evaluation focused on comparing automated OCR entity extraction against manual ground truth.

- **High Performance:** ``VITAL_NAME`` and ``TEST`` values showed >95% accuracy.
- **Low Performance:** ``Temporality`` and ``Assertion`` (specifically ``UNCERTAIN``) showed significant error rates (>25%).

2. Identified Systemic Weaknesses

1. **Temporality Drift:** The model frequently assigns ``CLINICAL_HISTORY`` to admission events that are clearly ``CURRENT``.
2. **Context Blindness:** Specific instructions (e.g., "Hold home meds") are often extracted as standard prescriptions without capturing the "Hold" instruction status accurately.

3. Proposed Guardrails

- **Keyword-Based Temporal Correction:** Implement a post-processing script to check for current dates associated with "admission" or "discharge" strings.
- **Negation/Hold Logic:** Introduce a dependency parser to link verbs like "Hold" or "Stop" directly to ``MEDICINE`` entities.